

Read Mapping and Variant Calling

Bioinformatics Workshop for *M. tuberculosis*
Genomics and Phylogenomics

July 10-14, 2018 @The Philippine Genome Center

Ulas Karaoz, PhD
Ecology Department
Berkeley Lab



Lawrence Berkeley
National Laboratory
Bringing Science Solutions to the World

Learning Objectives

- How read mappers work
- File formats involved in read mapping
- Querying and filtering read alignments
- Calling variants

What is Read Mapping?

genome

TGGCTTGC GGCCATCAACGGTT CCTACCGAGGGGGCGGTGCGGGC ATC
ACTAGTGC GGAA ACTGGGAGGGCTCTCGGCCCTCCGCCTT TAGGC GGTG
CTTACCCCTTCGTAGAGGGGCTGTGC GGCCATCCC CGC GAGGATCCGAGA
AGGCGAGCGTGC GGATCCCACCCGGGGGGACGGGCCTCAAAGCCGC TT

read

TGC GGCCATC

What is Read Mapping?

genome

TGGCTTGC~~GGCCATCAACGGTT~~CCTACCGAGGGGGCGGTGCGGGC~~ATC~~
ACTAGTGC~~GGAA~~ACTGGGAGGGCTCTCGGCC~~CT~~CCGC~~CTT~~AGGC~~GGT~~G
CTTACCC~~CT~~CGTAGAGGGCTGTGC~~GGCC~~ATCCC~~CG~~GAGGATCCGAGA
AGGCGAGCGTGC~~GG~~ATCCCACCCGGGGGACGGGC~~CT~~CAAAGCCGC~~TT~~

read

TGC~~GGCC~~ATC

read
mapping

Find all genomic locations "TGC~~GGCC~~ATC"
might have been generated from

TGGCTT**GC~~GGCC~~ATCAACGGTT**CCTACCGAGGGGGCGGT**GCGGGC**~~ATC~~
ACTAGTGC~~GGAA~~ACTGGGAGGGCTCTCGGCC~~CT~~CCGC~~CTT~~AGGC~~GGT~~G
CTTACCC~~CT~~CGTAGAGGGCTGT**GC~~GGCC~~ATCCC**~~CG~~GAGGATCCGAGA
AGGCGAGCGTGC~~GG~~ATCCCACCCGGGGGACGGGC~~CT~~CAAAGCCGC~~TT~~

What is Read Mapping?

genome

TGGCTTGC~~GGCC~~CATCAACGGTT~~CCTACCGAGGGGGCGGTGCGGGC~~ATC
ACTAGTGC~~GGAA~~ACTGGGAGGGCTCTCGGCC~~CTCCGC~~CTTAGGC~~GGTG~~
CTTACCC~~CTTC~~GTAGAGGGCTGTGC~~GGCC~~ATCCC~~CG~~GAGGATCCGAGA
AGGCGAGCGTGC~~GG~~ATCCCACC~~CGGGGG~~ACGGGC~~CTCAAAGCCGC~~CTT

read

TGC~~GGCC~~CATC

read
mapping

Find all genomic locations "TGC~~GGCC~~CATC"
might have been generated from

TGGCTT**TGC~~GGCC~~CATCAACGGTT**CCTACCGAGGGGGCGGTG**CGG**~~G~~CATC
ACTAGTGC~~GGAA~~ACTGGGAGGGCTCTCGGCC~~CTCCGC~~CTTAGGC~~GGTG~~
CTTACCC~~CTTC~~GTAGAGGGCTGT**GC~~GGCC~~ATCCC**~~CG~~GAGGATCCGAGA
AGGCGAGCGT**GC~~GG~~**ATCCCACC~~CGGGGG~~ACGGGC~~CTCAAAGCCGC~~CTT



- find all "approximate" occurrences
- analogous to string matching

Read Mapping isn't Read Alignment

read
alignment

TGGCTT**TGCGGCCAT**CAACGGTTCCCTACCGAGGGGGCGGT**TGCGGCCATC**
||||||| |||||
TGCGGCCATC ||||| | |||
TGCGGCCATC

ACTAGTGCGGAAACTGGGAGGGCTCTCGGCCCTCCGCCTTAGGCGGTG

CTTACCCCTTCGTAGAGGGGCT**TGCGGCCATC**CCCGCGAGGATCCGAGA
||||||| |||||
TGCGGCCATC

AGGCGAGCGT**TGCGG**--**AT**CCCACCCGGGGGACGGGCCTCAAAGCCGCCTT
||||| | |||
TGCGGCCATC

Read Mapping isn't Read Alignment

read
alignment

TGGCTT**TGCGGCCAT**CAACGGTTCCCTACCGAGGGGGCGGT**TGCGGCCATC**
||||||| |||||
TGCGGCCATC ||||| | |||
TGCGGCCATC

ACTAGTGCGGAAACTGGGAGGGCTCTCGGCCCTCCGCCTTAGGCGGTG

CTTACCCCTTCGTAGAGGGGCT**TGCGGCCATC**CCCGCGAGGATCCGAGA
||||||| |||||
TGCGGCCATC

AGGCGAGCGT**GCGG**--**AT**CCCACCCGGGGGACGGGCCTCAAAGCCGCCTT
||||| | |||
TGCGGCCATC

sequence alignment ≠ read mapping ≠ read alignment

Placing Read Mappers into the bigger picture

- Sensitive global aligners:
 - Needleman-Wuncsh
 - Smith-Waterman
 - Pair-HMM
- Pairwise heuristic:
 - fasta
 - Blast
 - Blat
 - Exonerate
- Whole genome:
 - Mavid
 - Mummer
 - Mauve
 - Lagan
 - BlastZ
- Short read:
 - Maq
 - SHRIMP
 - ELAND
 - bowtie
 - bwa
 - bbmap
 - SOAP

Read Mapping: Two Algorithmic Approaches

Read Mapping

I. Hash indexing

II. Burrows-Wheeler transform

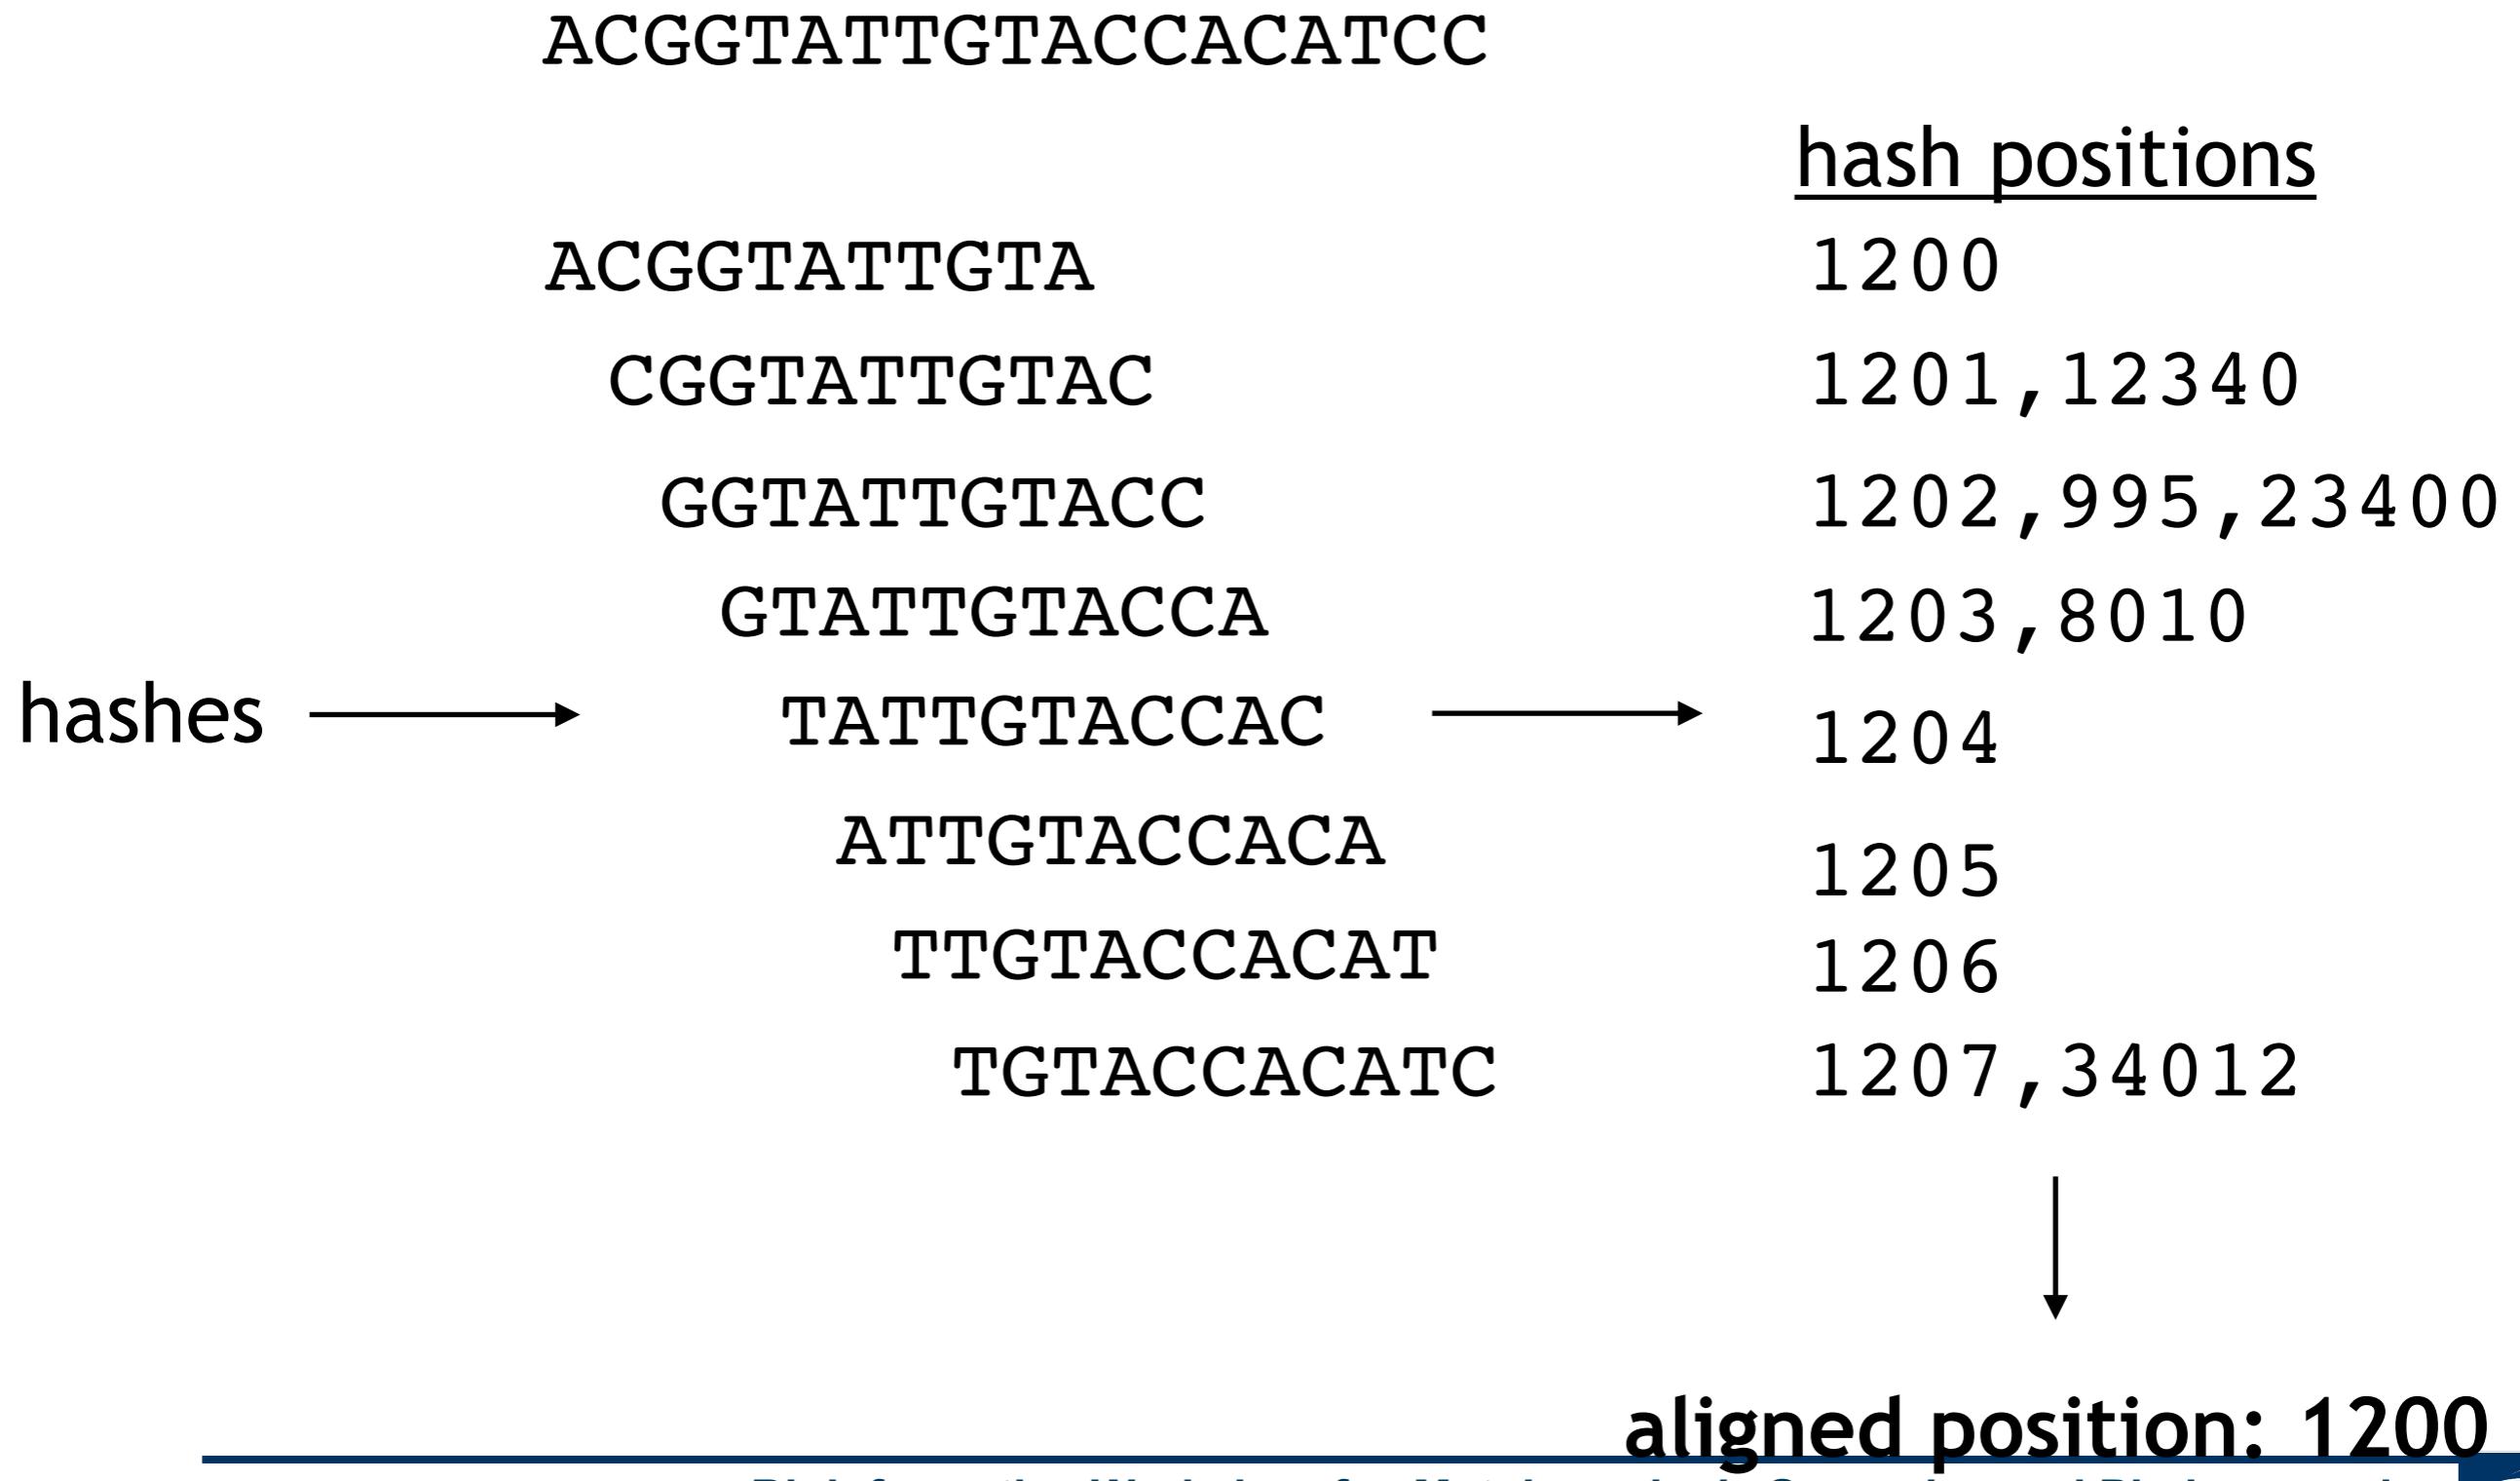
Read Mapping Approaches: I. Hash Indexing

Hash all genome k -mers using a hash table

key: all k -mers of fixed length k

value: genomic positions

Mapping Algorithms: Hashing



Read Mapping Approaches: II. Burrows-Wheeler Transform

ACAACG\$	ACAACG\$ \$ACAACG G\$ACAAC CG\$ACAA ACG\$ACA AACG\$AC CAACG\$A	\$ ACAACG AACG\$AC ACAACG\$ ACG\$ACA CAACG\$A CG\$ACAA G\$ACAA	GC\$AAAC
all rotations		lexical sorting	

Observe: rotations are *exposing* the suffixes

all suffixes	ACAACG\$ \$ G\$ CG\$ ACG\$ AACG\$ CAACG\$ ACAACG\$	\$ ACAACG AACG\$AC ACAACG\$ ACG\$ACA CAACG\$A CG\$ACAA G\$ACAA
--------------	-------------------------------------------------------------------------	----------------------------------------------------------------------------------

How all this is useful???

Burrows M, Wheeler DJ, A block sorting lossless data compression algorithm. Digital Equipment.

Read Mapping Approaches: II. Burrows-Wheeler Transform

\$ A C A A C G
A A C G \$ A C
A C A A C G \$
A C G \$ A C A
C A A C G \$ A
C G \$ A C A A
G \$ A C A A C

\$ A C A A C **G**₁
A₁ A C G \$ A **C**₁
A₂ C A A C G **\$**₁
A₃ C G \$ A C **A**₁
C₁ A A C G \$ **A**₂
C₂ G \$ A C A **A**₃
G₁ \$ A C A A **C**₂

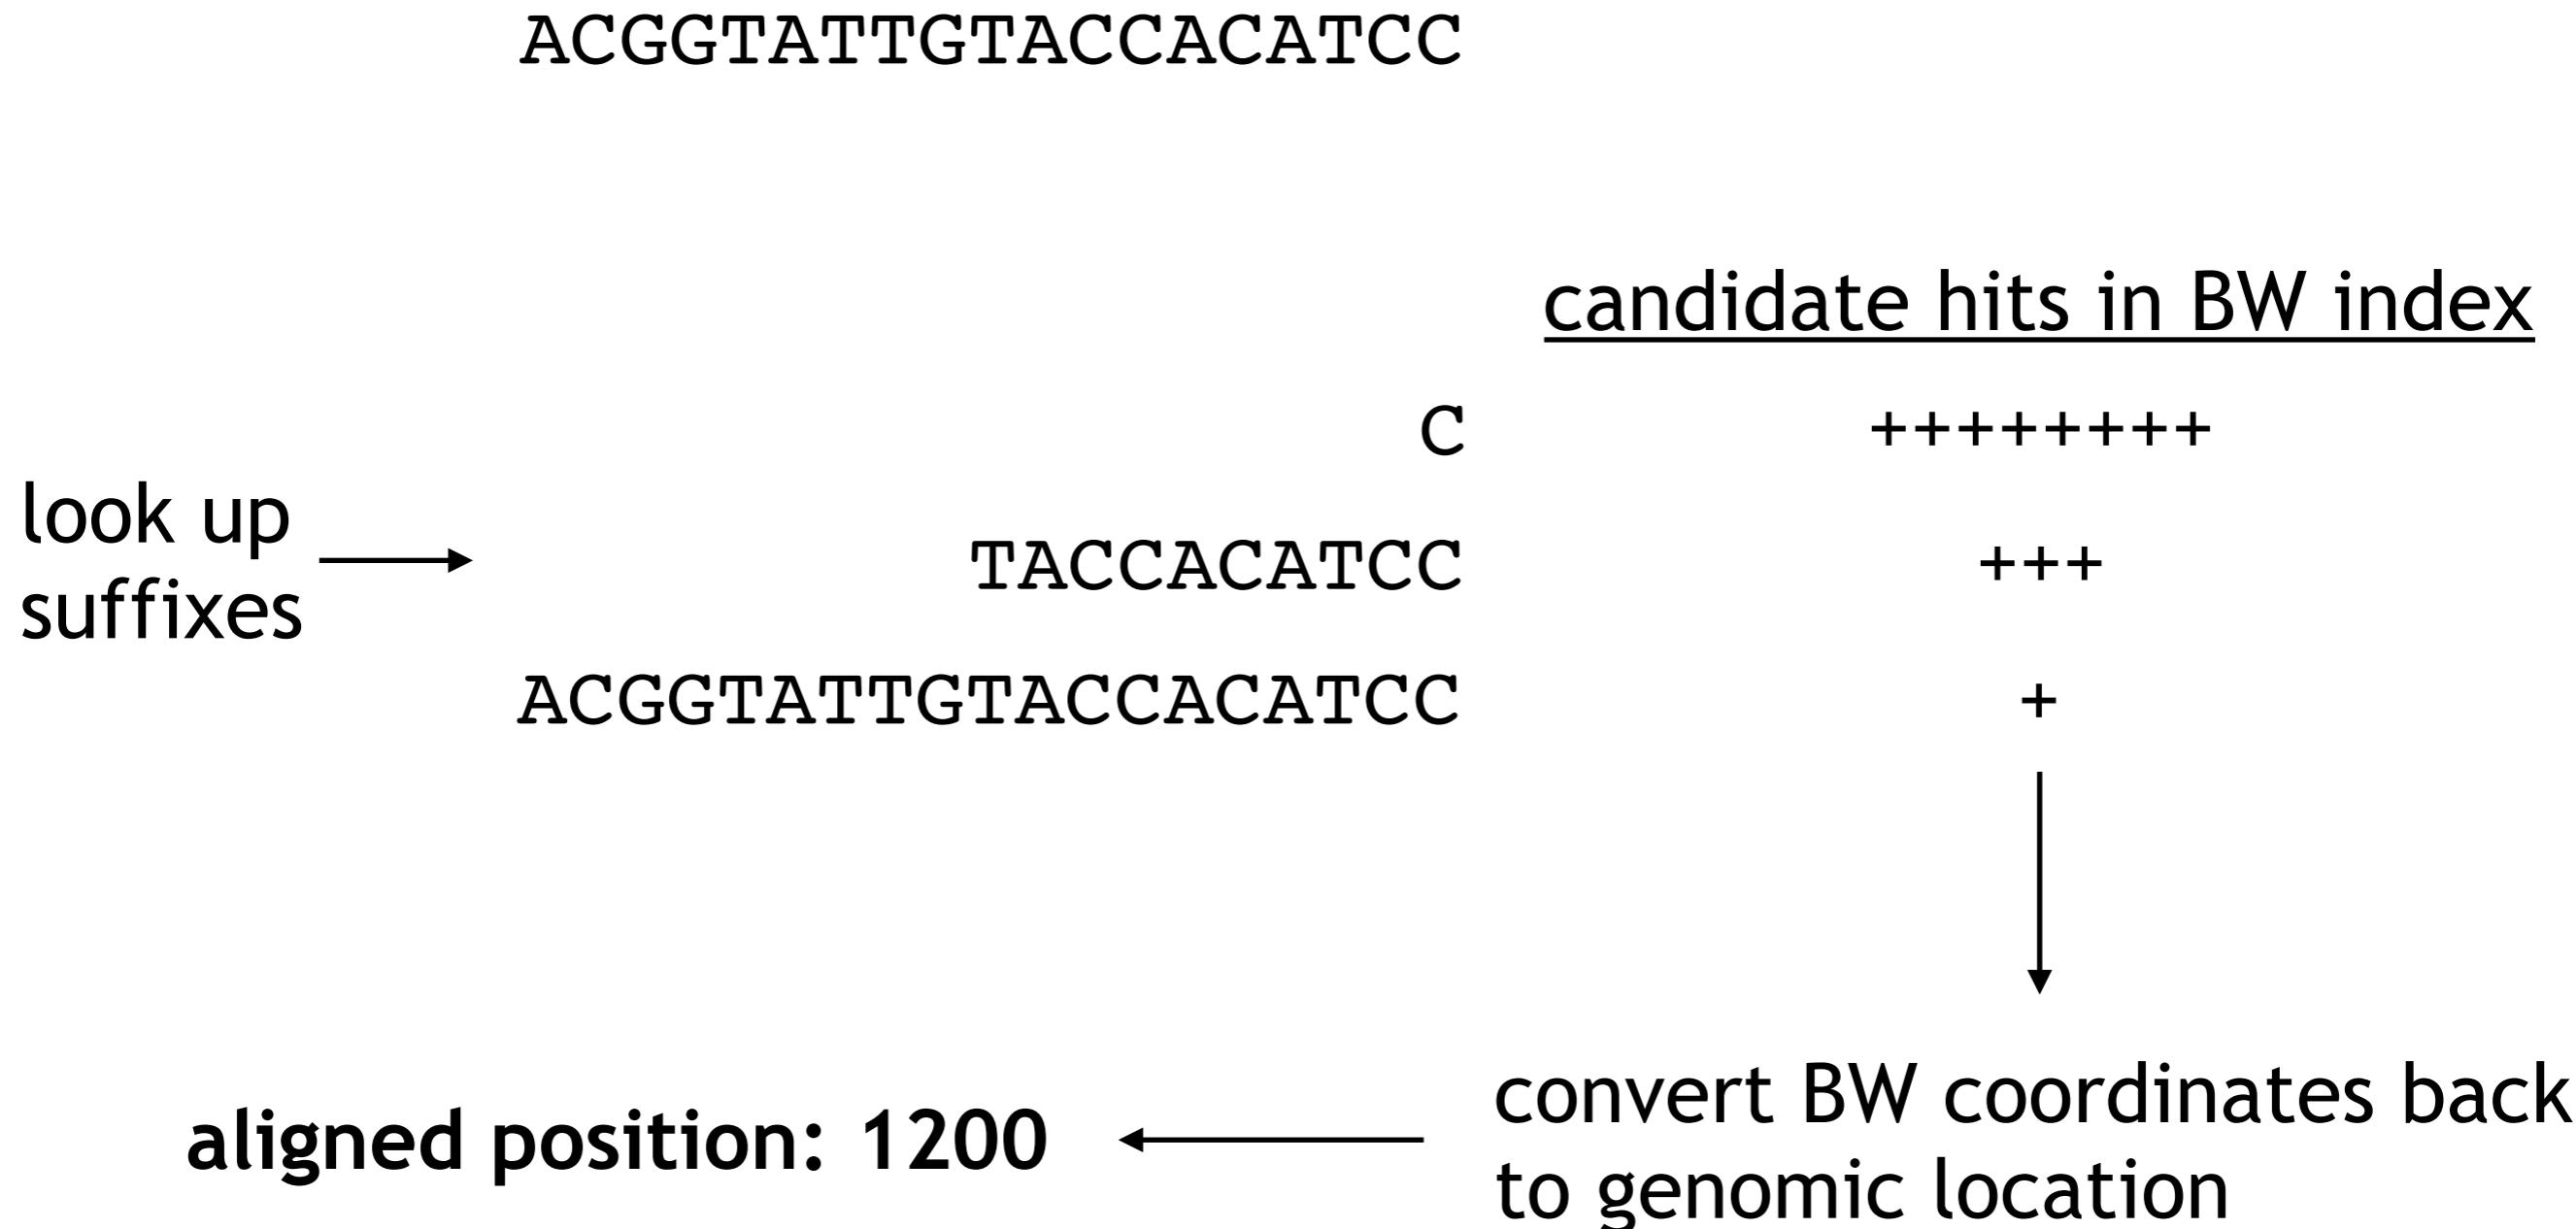
Can we reconstruct the original string from BWT?

$$\text{LF}(6, "C") = \text{occ}("C") + \text{count}(6, "C")$$

$$\text{LF}(6, "C") = \text{occ}("G") + \text{count}(6, "G")$$

Mapping Algorithms: BWT/Suffix Array

BWT: Burrows-Wheeler Transform



Mapping Quality

mapping quality: confidence in a read's genomic origin

- Probability of misalignment depends on:
 - uniqueness of the aligned region in the genome
 - alignment length
 - number of mismatches and gaps
- Expressed in Phred scores, similar to base qualities:
 - $AQ = -10 \log_{10}(P_{\text{misaligned}})$
- Values not standardized across different aligners. For instance, the value for "unique" mapping is not uniform across aligners.

Output Formats: SAM and BAM

SAM: Sequence Alignment/Map

optional header section

alignment section: one line per read,
containing 11 mandatory fields, followed by optional tags

```
@HD VN:1.4 SO:coordinate
@SQ SN:chr1 LN:249250621
@SQ SN:chr10 LN:135534747
@SQ SN:chr11 LN:135006516
@SQ SN:chr12 LN:133851895
@SQ SN:chr13 LN:115169878
@SQ SN:chr14 LN:107349540
@SQ SN:chr15 LN:102531392
@SQ SN:chr16 LN:90354753
@SQ SN:chr17 LN:81195210
@SQ SN:chr18 LN:78077248
@SQ SN:chr19 LN:59128983
@SQ SN:chr2 LN:243199373
@SQ SN:chr20 LN:63025520
@SQ SN:chr21 LN:48129895
@SQ SN:chr22 LN:51304566
@SQ SN:chr3 LN:198022430
@SQ SN:chr4 LN:191154276
@SQ SN:chr5 LN:180915260
@SQ SN:chr6 LN:171115067
@SQ SN:chr7 LN:159138663
@SQ SN:chr8 LN:146364022
@SQ SN:chr9 LN:141213431
@SQ SN:chrM LN:16571
@SQ SN:chrX LN:155270560
@SQ SN:chrY LN:59373566
@PG ID:TopHat VN:2.0.9 CL:/opt/chipster/tools/tophat2/tophat -p 2 --read-mismatches 2 -a 8 -m 0 -i 70 -I 500000 -g 20 --library-type fr-unstranded
--transcriptome-index=/opt/chipster/tools/bowtie2/indexes/hg19.ti --no-novel-juncs /opt/chipster/tools/bowtie2/indexes/hg19 reads1.fq
HWI-EAS229_1:4:82:1371:1147 272 chr1 18378 1 2M6358N73M* 0 0
TCCTGCTGAAGATGTCTCCAGAGACCTCTGCAGGTACTGAAGGGCATCCGCCATCTGCTGGACGGCCTCCTCTC 5661525416816488666(6(6(6261?8==(B=513);(/BB=141=>6>?==<?B>9B?>BA<66>BA>BB
CC:Z:chr15MD:Z:40C34XG:i:0 NH:i:3 HI:i:0 NM:i:1 XM:i:1 XN:i:0 XO:i:0 CP:i:102506354 AS:i:0 XS:A:- YT:Z:UU
```

BAM: binary version of SAM

Output Formats: SAM and BAM

- The specification
 - <http://samtools.sourceforge.net/SAM1.pdf>
- The SAM format consists of two sections:
 - Header section: Used to describe source of data, reference sequence, method of alignment, etc.
 - Alignment section: Used to describe the read, quality of the read, and nature alignment of the read to a region of the genome
- BAM is a compressed version of SAM
 - Compressed using lossless BGZF format
 - Other BAM compression strategies are a subject of research. See ‘CRAM’ format for example
- BAM files are usually ‘indexed’
 - A ‘.bai’ file will be found beside the ‘.bam’ file
 - Indexing aims to achieve fast retrieval of alignments overlapping a specified region without going through the whole alignments. BAM must be sorted by the reference ID and then the leftmost coordinate before indexing

SAM/BAM header section

- Used to describe source of data, reference sequence, method of alignment, etc.
- Each section begins with character '@' followed by a two-letter record type code. These are followed by two-letter tags and values
 - @HD The header line
 - VN: format version
 - SO: Sorting order of alignments
 - @SQ Reference sequence dictionary
 - SN: reference sequence name
 - LN: reference sequence length
 - SP: species
 - @RG Read group
 - ID: read group identifier
 - CN: name of sequencing center
 - SM: sample name
 - @PG Program
 - PN: program name
 - VN: program version

<https://samtools.github.io/hts-specs/SAMv1.pdf>

Genomics and Phylogenomics

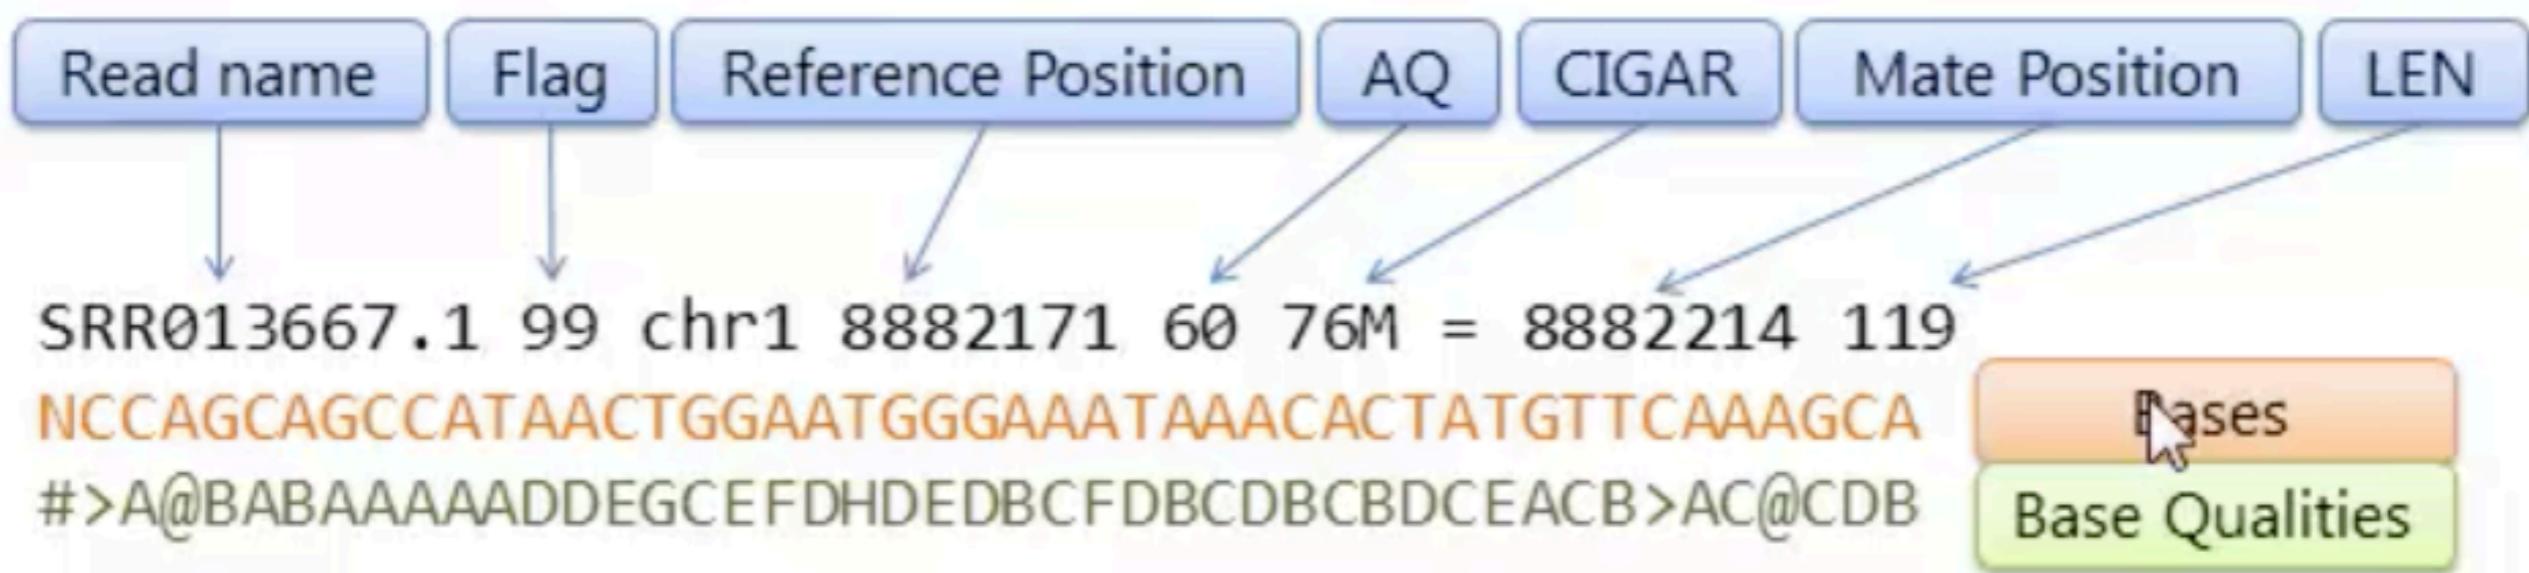
@The Philippine Genome Center



BERKELEY LAB

SAM/BAM alignment section

Col	Field	Type	Regexp/Range	Brief description
1	QNAME	String	[!-?A-~]{1,255}	Query template NAME
★ 2	FLAG	Int	[0,2 ¹⁶ -1]	bitwise FLAG
3	RNAME	String	* [!-()+-<>-~] [!-~]*	Reference sequence NAME
4	POS	Int	[0,2 ²⁹ -1]	1-based leftmost mapping POsition
5	MAPQ	Int	[0,2 ⁸ -1]	MAPping Quality
★ 6	CIGAR	String	* ([0-9]+[MIDNSHPX=])+	CIGAR string
7	RNEXT	String	* = [!-()+-<>-~] [!-~]*	Ref. name of the mate/next segment
8	PNEXT	Int	[0,2 ²⁹ -1]	Position of the mate/next segment
9	TLEN	Int	[-2 ²⁹ +1,2 ²⁹ -1]	observed Template LENgth
10	SEQ	String	* [A-Za-z.=.]+	segment SEQuence
11	QUAL	String	[!-~]+	ASCII of Phred-scaled base QUALity+33



SAM/BAM flag

- <http://broadinstitute.github.io/picard/explain-flags.html>
- 12 bitwise flags describing the alignment
- These flags are stored as a binary string of length 11 instead of 11 columns of data
- Value of '1' indicates the flag is set. e.g. 00100000000
- All combinations can be represented as a number from 1 to 2048 (i.e. $2^{11}-1$). This number is used in the BAM/SAM file. You can specify 'required' or 'filter' flags in samtools view using the '-f' and '-F' options respectively

Bit	Description
1	0x1 template having multiple segments in sequencing
2	0x2 each segment properly aligned according to the aligner
4	0x4 segment unmapped
8	0x8 next segment in the template unmapped
16	0x10 SEQ being reverse complemented
32	0x20 SEQ of the next segment in the template being reverse complemented
64	0x40 the first segment in the template
128	0x80 the last segment in the template
256	0x100 secondary alignment
512	0x200 not passing filters, such as platform/vendor quality controls
1024	0x400 PCR or optical duplicate
2048	0x800 supplementary alignment

Note that to maximize confusion, each bit is described in the SAM specification using its hexadecimal representation (i.e., '0x10' = 16 and '0x40' = 64).

<https://samtools.github.io/hts-specs/SAMv1.pdf>

Genomics and Phylogenomics

@The Philippine Genome Center

CIGAR String

Op	BAM	Description
M	0	alignment match (can be a sequence match or mismatch)
I	1	insertion to the reference
D	2	deletion from the reference
N	3	skipped region from the reference
S	4	soft clipping (clipped sequences present in SEQ)
H	5	hard clipping (clipped sequences NOT present in SEQ)
P	6	padding (silent deletion from padded reference)
=	7	sequence match
X	8	sequence mismatch

- The CIGAR string is a sequence of base lengths and associated ‘operations’ that are used to indicate which bases align to the reference (either a match or mismatch), are deleted, are inserted, represent introns, etc.
- e.g. 81M859N19M

BED format: specify genomic regions as coordinates

- When working with BAM files, it is very common to want to examine a focused subset of the reference genome
- These subsets are commonly specified in ‘BED’ files
 - <https://genome.ucsc.edu/FAQ/FAQformat.html#format1>
- Many BAM manipulation tools accept regions of interest in BED format
- Basic BED format (tab separated):
 - Chromosome name, start position, end position
 - Coordinates in BED format are 0 based

Manipulation of BAM, SAM and BED files

- Several tools are used ubiquitously to query and manipulate these files
- SAM/BAM files
 - samtools
 - bamtools
 - picard
- BED files
 - bedtools
 - bedops

Sorting SAM/BAM

- Generally, BAM files are sorted by position
 - For computational performance reasons, when sorted and indexes, arbitrary access is much faster
- Certain tools require a BAM sorted by "*read name*"
 - Usually, when you need to easily identify both reads of a pair
 - Example: check if the insert size looks right

Alignment QC

Alignment Level QC

- How many reads mapped to the reference?
 - How many mapped uniquely?
- How many pairs mapped?
 - How many pairs mapped concordantly?
- Mapping quality distribution
- Quick and dirty (*samtools flagstat*)

52841623 + 0 in total (QC-passed reads + QC-failed reads)

0 + 0 duplicates

52841623 + 0 mapped (100.00%:-nan%)

52841623 + 0 paired in sequencing

28919461 + 0 read1

23922162 + 0 read2

42664064 + 0 properly paired (80.74%:-nan%)

44904884 + 0 with itself and mate mapped

7936739 + 0 singletons (15.02%:-nan%)

999152 + 0 with mate mapped to a different chr

357082 + 0 with mate mapped to a different chr (mapQ>=5)

- Detailed: *qualimap*

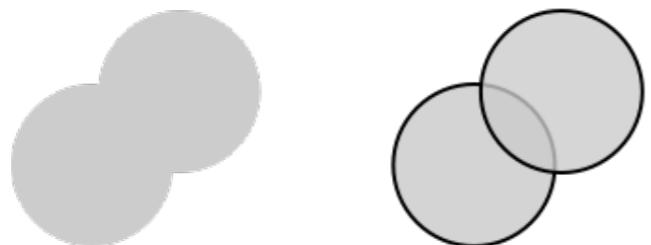
Duplicate reads (fragments)

Duplicate reads (fragments)

- If detected, essential to remove them before variant calling
- **Optical duplicates** (hint: think how Illumina sequencing by synthesis works)
 - generated when a single cluster of reads is part of two adjacent tiles' on the same slide and used to compute two read calls separately
 - very similar in sequence (except sequencing errors)
 - identified where the first, say, 50 bases are identical between two reads and the read's coordinates are close
- **Library duplicates** (only relevant for targeted sequencing where there is pre-library PCR)
 - generated when the original sample is preamplified to such extent that initial unique targets are PCR replicated prior to library preparation and will lead to several independent spots on the Illumina slide.
 - do not have to be adjacent on the slide
 - share a very high level of sequence identity
 - align to the same place on reference
 - identified from alignment to reference

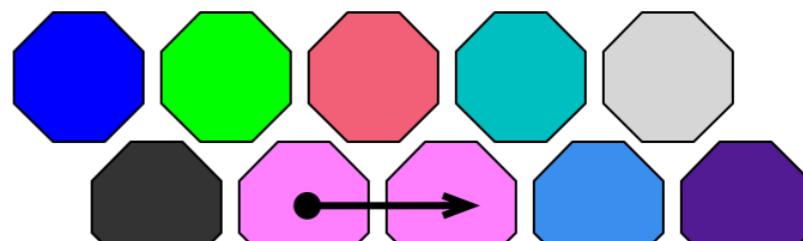
Duplicate Reads (Illumina)

"Optical" duplicates: caused by Illumina's optical detection/image processing



- A single cluster that has falsely been called as two by RTA
- single cluster
- image processing calls 2 clusters
- doesn't happen on patterned flow cells

"Exclusion Amplification (ExAmp)" duplicates:
During cluster generation, a single cluster might occupy two adjacent wells.

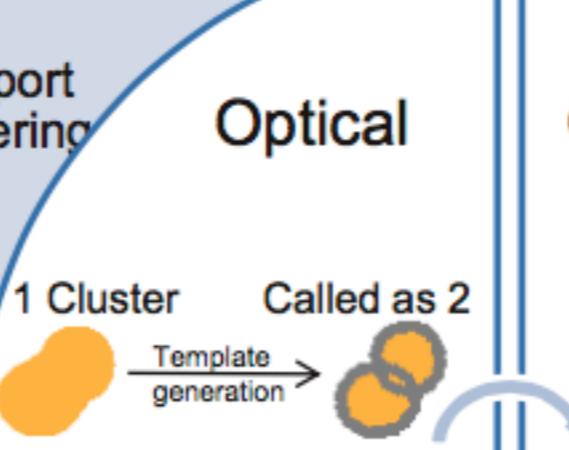


- Wells on MiSeq 3000/4000
- only on patterned flow cells

"PCR/enrichment" duplicates:

- During library construction fragments with adapters are enriched by PCR amplification.

Not on Patterned Flow Cells



Clustering

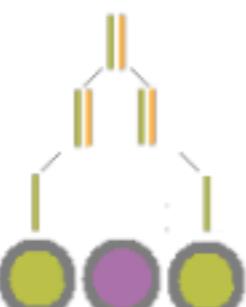


- During cluster generation a library occupies two adjacent wells

Unique to Patterned Flow Cells

- Duplicate molecules that arise from amplification
- during sample prep

PCR



Sister



Complement strands of same library form independent clusters

- Treated as duplicates by some informatic pipelines

Mappers are less accurate around indel sites

Solution: do local realignment post mapping

Reference CTTTAGTTCTTT----CTTTCTTCTTCCTTTTAAGTCTCCCTC

Reads CTTTAGTTCTTT----GCCGCTTCTTCCTTCTT

CTTTAGTTCTTT----GCCGCTTCTTCCTTCTT

Reads CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

Reference CTTTAGTTCTTT----CTTTCTTCTTCCTTTCTTCTTCTTCTT

Reads CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTT

Reads CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

CTTTAGTTCTTTGCCGCTTCTTCCTTCTTCTTCTTCTTCTTCTTCTTCTT

Ambiguous alignment around adjacent SNPs

Reads	Reference	Reads
	AAGCGTCG	AAG---CGTCG
	AAGCGTCG	AAG---CGTCG
	AAGCGTCG	AAG---CGTCG
	AAGCTACG	AAGCTACG
	AAGCTACG	AAGCTACG
	AAGCTACG	AAGCTACG

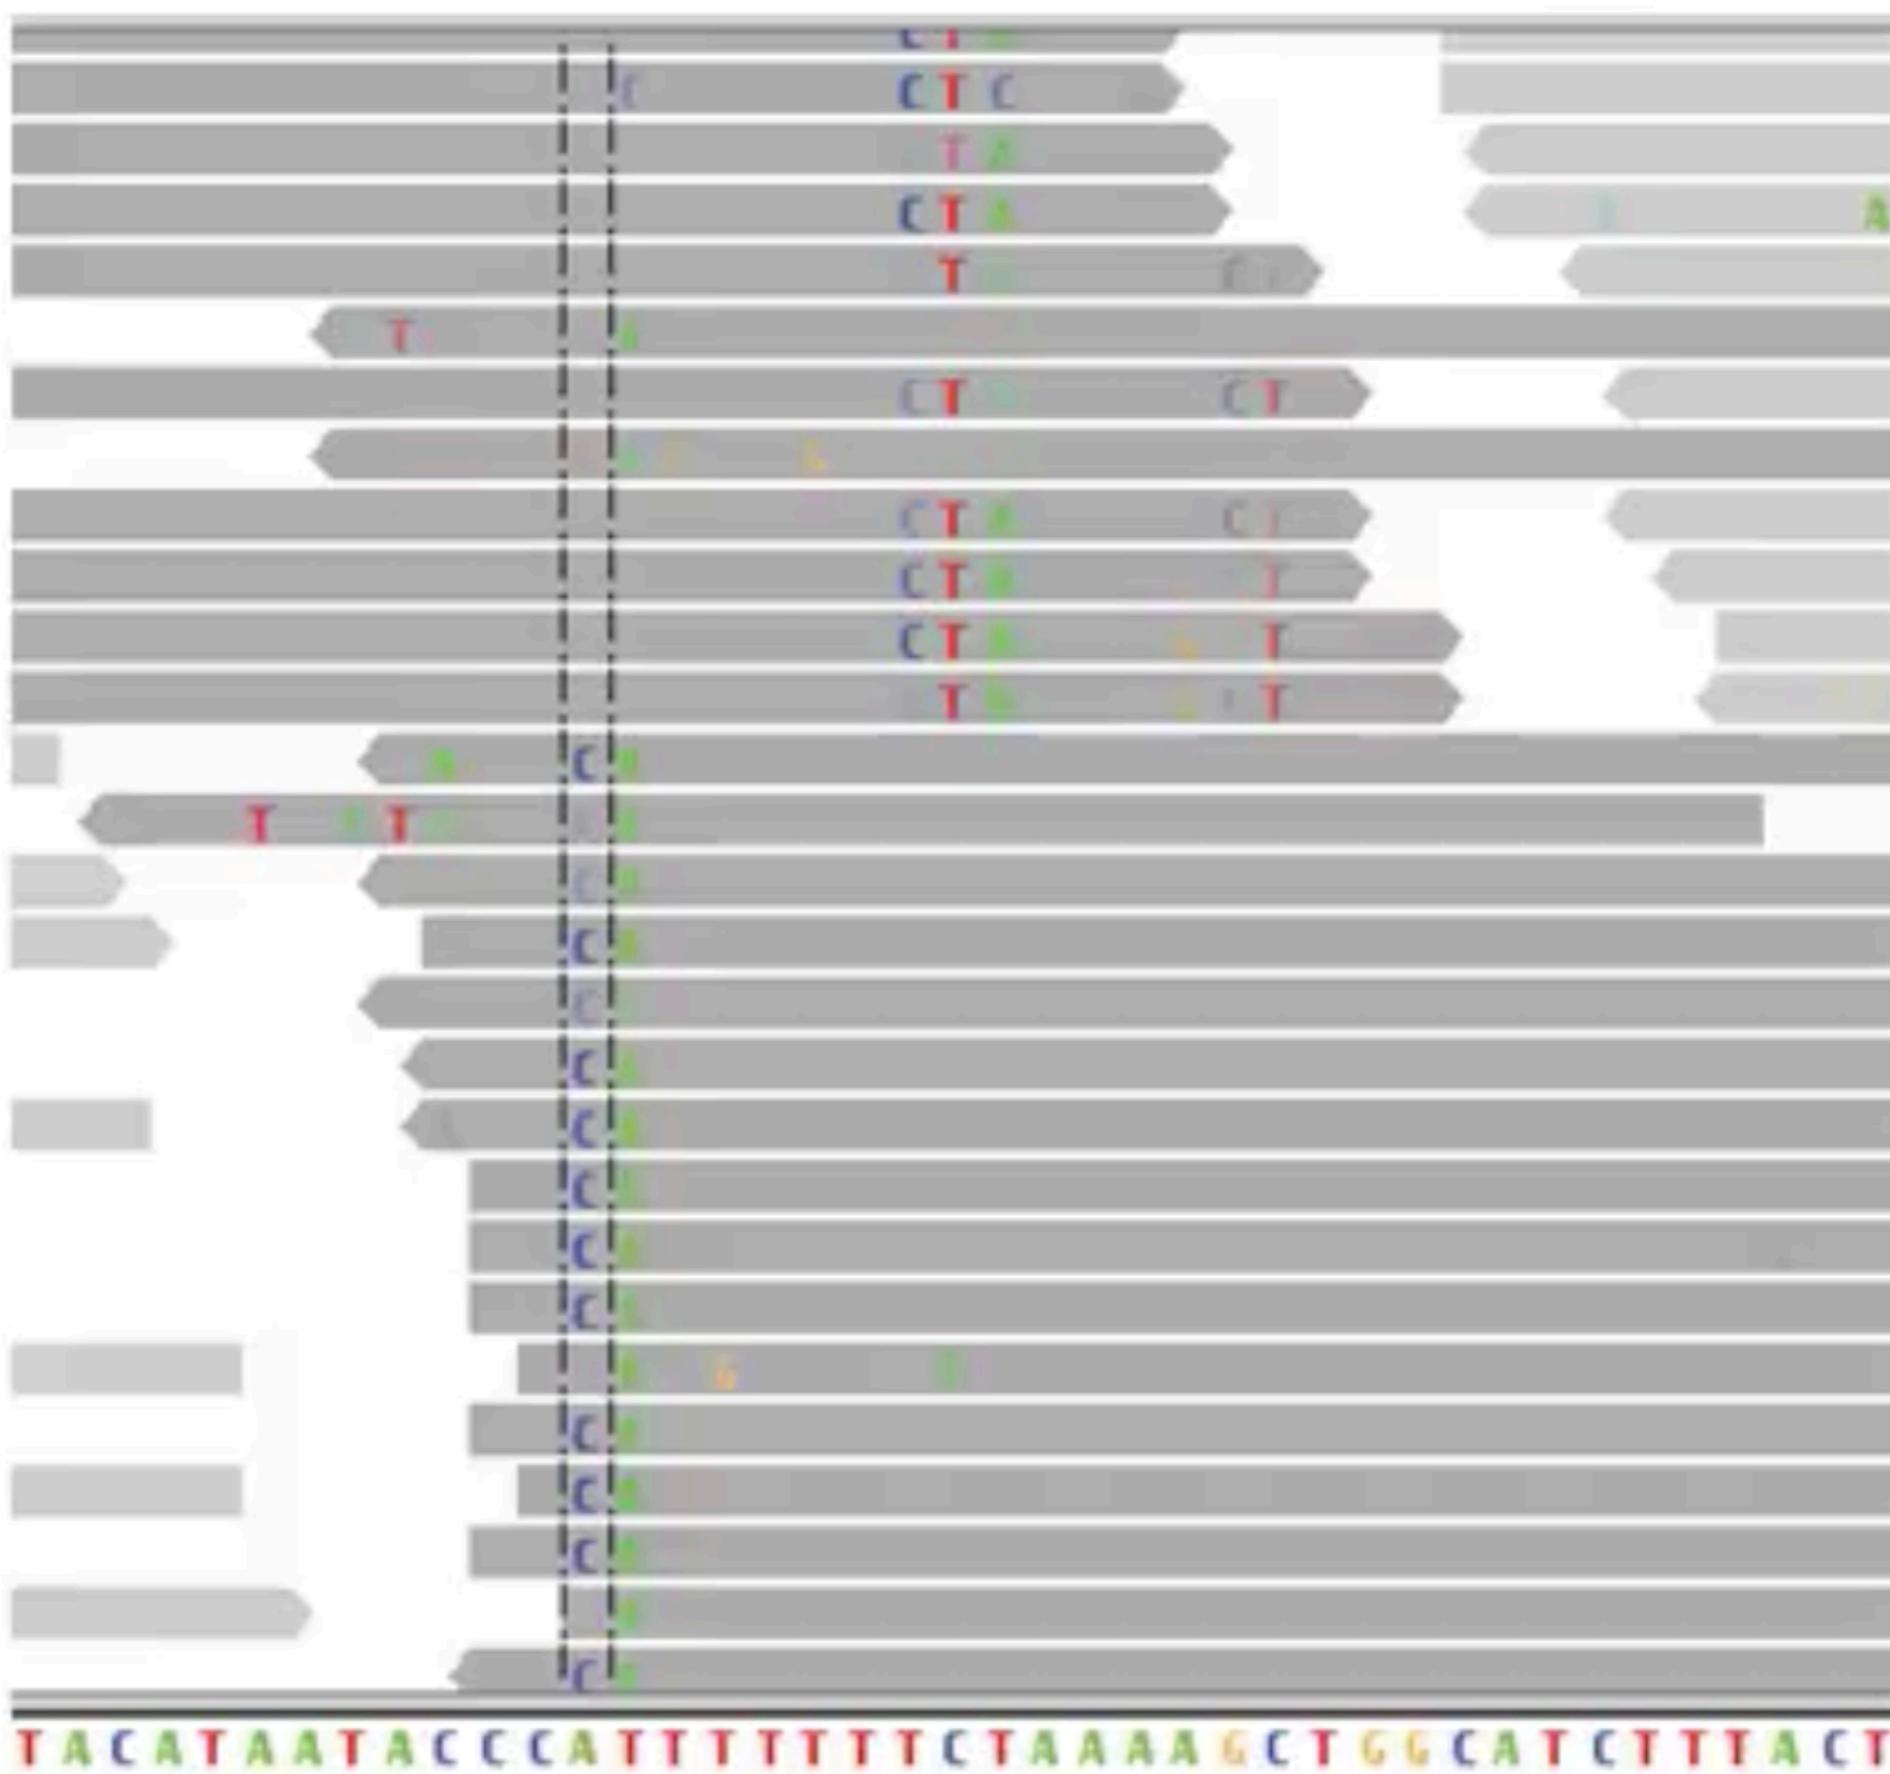
3 adjacent SNPs OR an insertion

Ambiguous alignment homopolymer runs flanked by adjacent SNPs

Reads

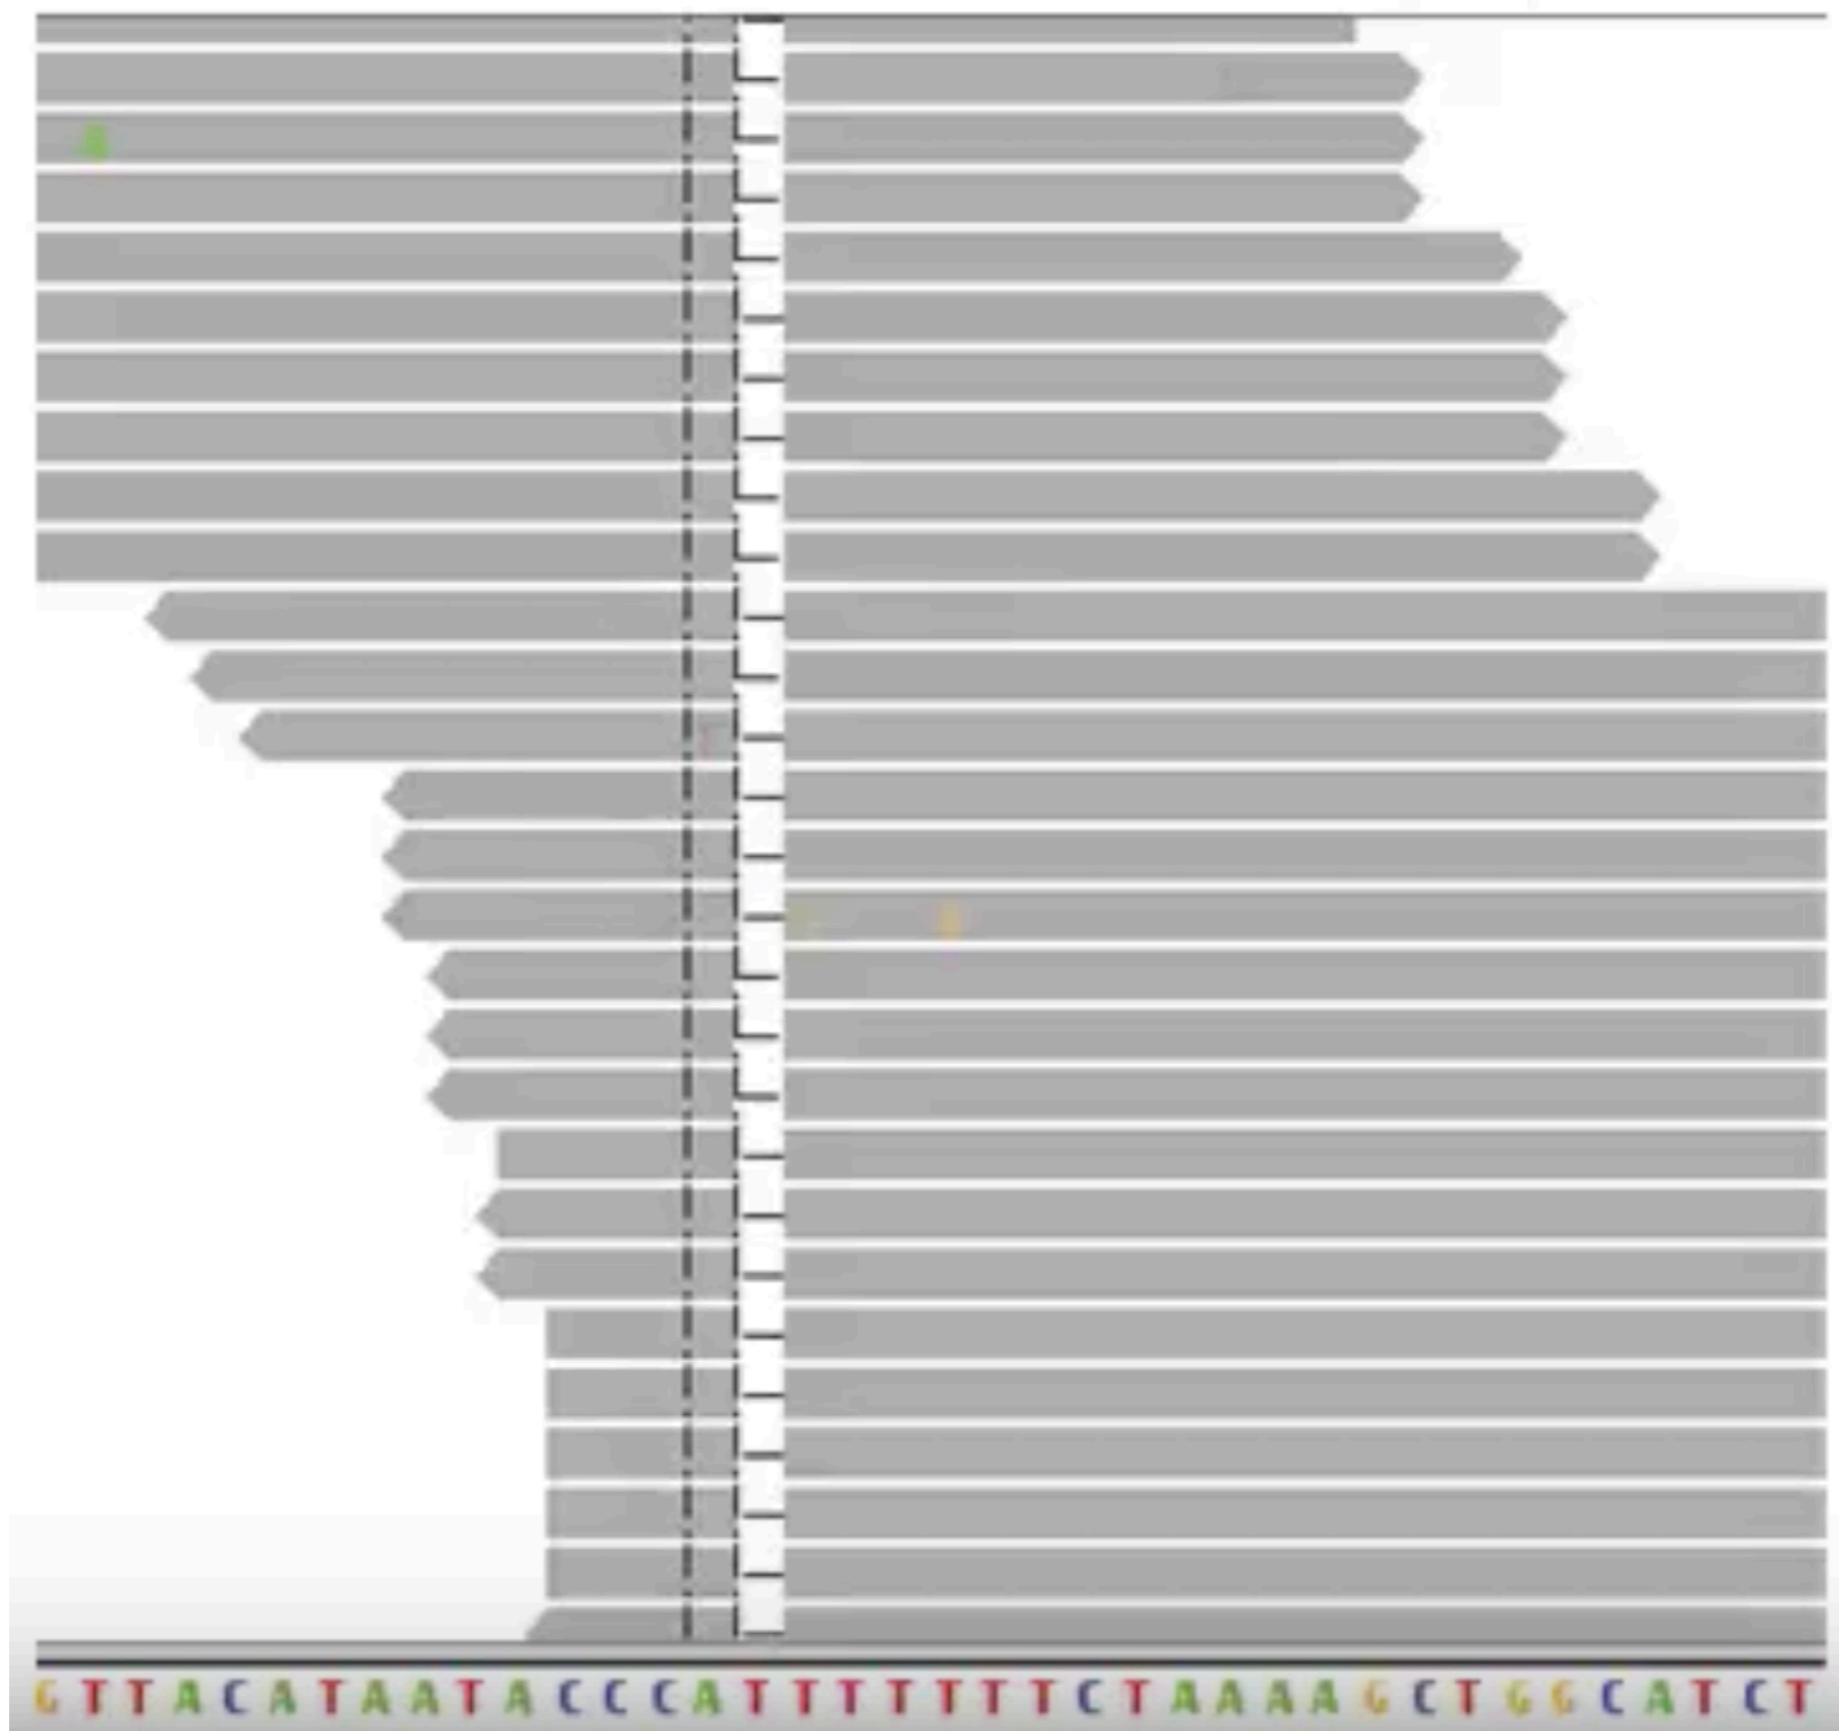
Reference	Reads	Reference	Reads
... CCCATTTCCTAAAAGCTGGCAT CCCATTTCCTAAAAGCTGGCAT ...	
CCC A TTTCCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...
CCC A TTTCCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...
CCC A TTTCCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...	CCC A -TTTCTAAAAGCTGGCAT ...
... CCCATTTC T AAAAAA	... CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT ...
... CCCATTTC T AAAAAA	... CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT ...
... CCCATTTC T AAAAAA	... CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT CCCA-TTTCTAAAAGCTGGCAT ...

Alignment post-processing: Indel realignment



Alignment post-processing: Indel realignment

Alignment post-processing: Indel realignment



What makes a good aligner?

- Speed
- Accuracy
 - Novoalign
 - Razers3
- All-round
 - bwa/bwa-mem
 - bowtie
- Functionality
 - STAR
 - TopHat

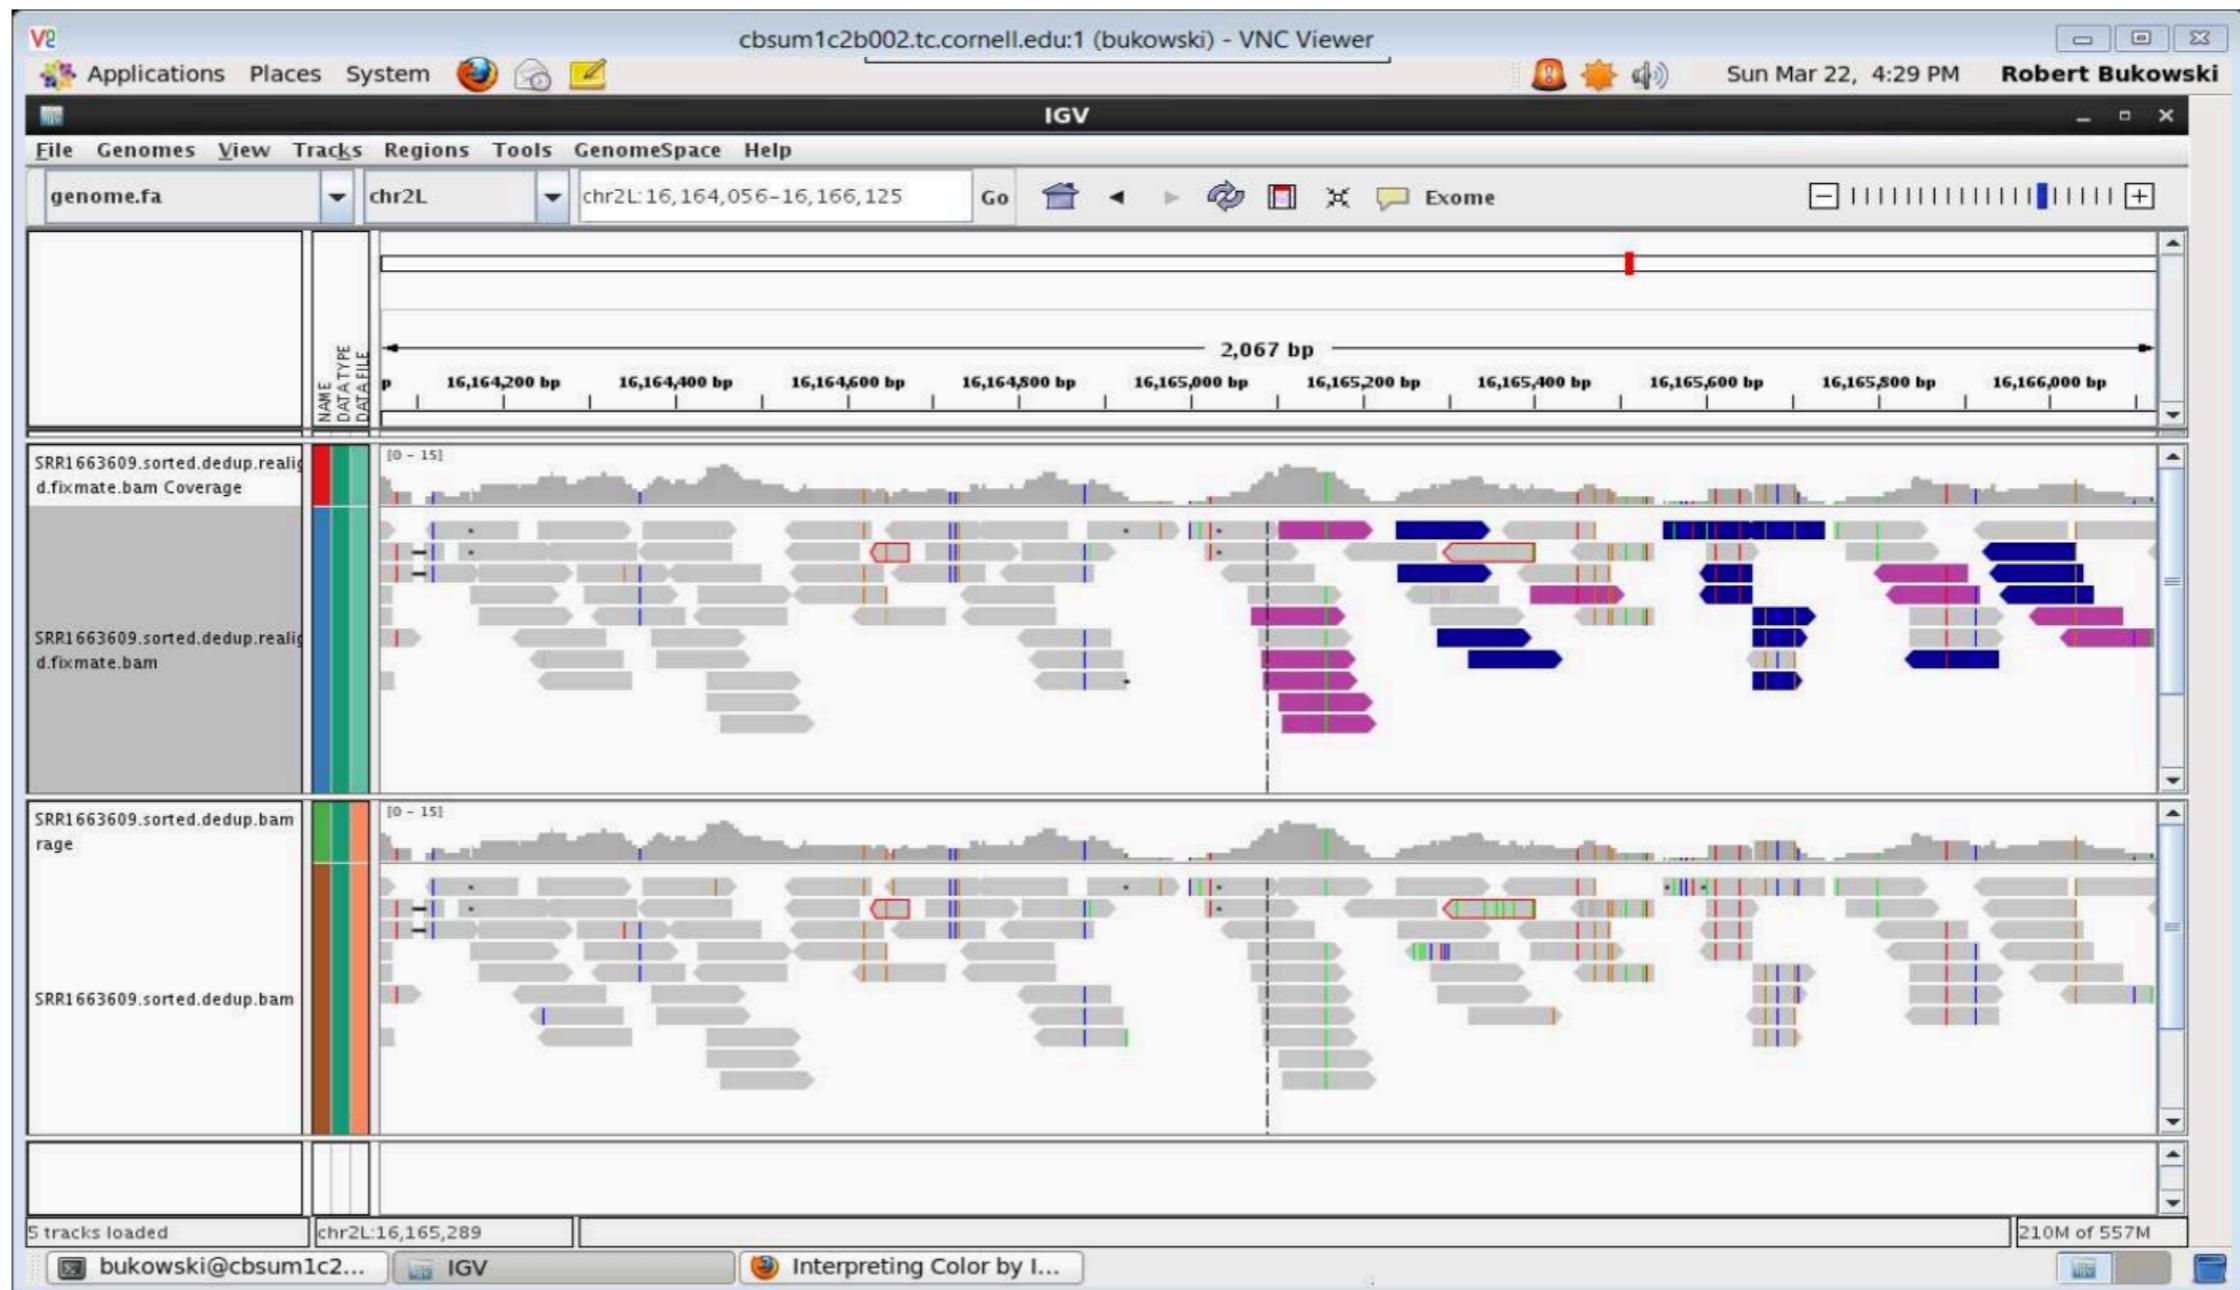
Visualizing read mapping/alignment: IGV

Look at multiple BAM files

Zoom in and out

Various color-coding schemes

Can load genome annotation track



VCF: Variant Call Format

Describes information about sequence variation.

BCF is the compressed version.

```
##fileformat=VCFv4.1
##fileDate=20180619
##source="Pilon version 1.22 Wed Mar 15 16:38:30 2017 -0400"
##PILON="--genome /Users/ukaraoz/Work/MTB/SFClusters/bin/data/ref/GCA_000277735.2_ASM27773v2_genomic.fna --bam test.sorted.duprem.bam --output out_pilon --variant"
##reference=file:/Users/ukaraoz/Work/MTB/SFClusters/bin/data/ref/GCA_000277735.2_ASM27773v2_genomic.fna
##contig=<ID=CP003248.2,length=4411709>
##FILTER=<ID=LowCov,Description="Low Coverage of good reads at location">
##FILTER=<ID=Amb,Description="Ambiguous evidence in haploid genome">
##FILTER=<ID=Del,Description="This base is in a deletion or change event from another record">
##INFO=<ID=DP,Number=1,Type=Integer,Description="Valid read depth; some reads may have been filtered">
##INFO=<ID=TD,Number=1,Type=Integer,Description="Total read depth including bad pairs">
##INFO=<ID=PC,Number=1,Type=Integer,Description="Physical coverage of valid inserts across locus">
##INFO=<ID=BQ,Number=1,Type=Integer,Description="Mean base quality at locus">
##INFO=<ID=MQ,Number=1,Type=Integer,Description="Mean read mapping quality at locus">
##INFO=<ID=QD,Number=1,Type=Integer,Description="Variant confidence/quality by depth">
##INFO=<ID=BC,Number=4,Type=Integer,Description="Count of As, Cs, Gs, Ts at locus">
##INFO=<ID=QP,Number=4,Type=Integer,Description="Percentage of As, Cs, Gs, Ts weighted by Q & MQ at locus">
##INFO=<ID=IC,Number=1,Type=Integer,Description="Number of reads with insertion here">
##INFO=<ID=DC,Number=1,Type=Integer,Description="Number of reads with deletion here">
##INFO=<ID=XC,Number=1,Type=Integer,Description="Number of reads clipped here">
##INFO=<ID=AC,Number=A,Type=Integer,Description="Allele count in genotypes, for each ALT allele, in the same order as listed">
##INFO=<ID=AF,Number=A,Type=Float,Description="Fraction of evidence in support of alternate allele(s)">
##INFO=<ID=SVTYPE,Number=1,Type=String,Description="Type of structural variant">
##INFO=<ID=SVLEN,Number=.,Type=String,Description="Difference in length between REF and ALT alleles">
##INFO=<ID=END,Number=1,Type=Integer,Description="End position of the variant described in this record">
##INFO=<ID=IMPRECISE,Number=0,Type=Flag,Description="Imprecise change from local reassembly (ALT contains Ns)">
##FORMAT=<ID=GT,Number=1,Type=String,Description="Genotype">
##FORMAT=<ID=AD,Number=.,Type=String,Description="Allelic depths for the ref and alt alleles in the order listed">
##FORMAT=<ID=DP,Number=1,Type=String,Description="Approximate read depth; some reads may have been filtered">
##ALT=<ID=DUP,Description="Possible segmental duplication">
#CHROM POS ID REF ALT QUAL FILTER INFO FORMAT SAMPLE
```

FILTER

INFO

FORMAT

VARIANTS

CP003248.2	1977	.	A	G	2357	PASS	DP=59;TD=59;BQ=40;MQ=60;QD=39;BC=0,0,59,0;QP=0,0,100,0;PC=90;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	2532	.	T	C	2298	PASS	DP=57;TD=57;BQ=40;MQ=60;QD=40;BC=0,57,0,0;QP=0,100,0,0;PC=85;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	4013	.	T	C	2629	PASS	DP=65;TD=65;BQ=40;MQ=60;QD=40;BC=0,65,0,0;QP=0,100,0,0;PC=95;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	6112	.	G	C	2726	PASS	DP=69;TD=69;BQ=40;MQ=60;QD=39;BC=0,69,0,0;QP=0,100,0,0;PC=101;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	6832	.	C	T	2427	PASS	DP=60;TD=60;BQ=40;MQ=60;QD=40;BC=0,0,0,60;QP=0,0,0,100;PC=88;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	7297	.	A	G	2674	PASS	DP=70;TD=70;BQ=39;MQ=60;QD=38;BC=0,1,69,0;QP=0,0,100,0;PC=101;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	24698	.	GGCGCGTTGCTCGGGTAA	G	.	PASS	DP=72;TD=73;BQ=40;MQ=59;QD=10;BC=0,19,0,0;QP=0,100,0,0;PC=106;IC=0;DC=53;XC=0;AC=2;AF=0.74	GT 1/1 1
CP003248.2	24716	.	A	G	371	Del;Amb	DP=15;TD=68;BQ=183;MQ=12;QD=24;BC=9,0,5,1;QP=58,0,35,6;PC=105;IC=0;DC=53;XC=0;AC=1;AF=0.38	GT 0/1 1
CP003248.2	24720	.	CGTTGCTCGGGTAACCGC	C	.	PASS	SVTYPE=DEL;SVLEN=-18;END=24738	GT 1/1 1
CP003248.2	71581	.	C	CCGAGCGCTGTTCTGGCGCTAACATCTGACGCTAGAATAG	2808	PASS	DP=71;TD=71;BQ=40;MQ=47;QD=39;BC=0,71,0,0;QP=0,100,0,0;PC=98;IC=17;DC=0;XC=0;AC=2;AF=0.30	GT 1/1 1
CP003248.2	150039	.	G	T	4	Amb;LowCov	DP=4;TD=4;BQ=40;MQ=60;QD=1;BC=0,0,3,1;QP=0,0,74,26;PC=28;IC=0;DC=0;XC=0;AC=1;AF=0.26	GT 0/1 1
CP003248.2	150893	.	G	GGGCCCGCGGTCAAGGCAGGGCGTCAGGCCTGGCGCTAACATATCGGCCGACCGGGGGCCGGCGAACGGCGGGCCGGC	.	PASS	SVTYPE=INS;SVLEN=181;END=150893;IMPRECISE	GT 1/1 1
CP003248.2	154189	.	A	G	844	Del	DP=21;TD=41;BQ=40;MQ=51;QD=40;BC=0,0,21,0;QP=0,0,100,0;PC=86;IC=0;DC=0;XC=2;AC=2;AF=1.00	GT 1/1 1
CP003248.2	1618984	.	T	C	160	LowCov	DP=4;TD=4;BQ=40;MQ=60;QD=40;BC=0,4,0,0;QP=0,100,0,0;PC=108;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT 1/1 1
CP003248.2	1637527	.	CCCCGCCGGG	C	.	Amb;LowCov	DP=5;TD=5;BQ=39;MQ=0;QD=31;BC=0,4,0,0;QP=0,100,0,0;PC=110;IC=0;DC=1;XC=0;AC=1;AF=0.20	GT 0/1 1
CP003248.2	3120174	.	T	C	1	Del;Amb;LowCov	DP=5;TD=23;BQ=41;MQ=3;QD=0;BC=0,1,0,4;QP=0,25,0,75;PC=102;IC=0;DC=0;XC=1;AC=1;AF=0.25	GT 0/1 1
CP003248.2	4386286	.	A	G	4633	Del;Amb	DP=158;TD=161;BQ=40;MQ=60;QD=29;BC=69,0,89,0;QP=43,0,57,0;PC=200;IC=0;DC=0;XC=0;AC=1;AF=0.57	GT 0/1 1

VCF: Variant Call Format

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	SAMPLE
CP003248.2	1977	.	A	G	2357	PASS	DP=59;TD=59;BQ=40;MQ=60;QD=39;BC=0,0,59,0;QP=0,0,100,0;PC=90;IC=0;DC=0;XC=0;AC=2;AF=1.00	GT	1/1

- DP Valid read depth; some reads may have been filtered">
TD Total read depth including bad pairs">
PC Physical coverage of valid inserts across locus">
BQ Mean base quality at locus">
MQ Mean read mapping quality at locus">
QD Variant confidence/quality by depth">
BC Count of As, Cs, Gs, Ts at locus">
QP Percentage of As, Cs, Gs, Ts weighted by Q & MQ at locus">
IC Number of reads with insertion here">
DC Number of reads with deletion here">
XC Number of reads clipped here">
AC Allele count in genotypes, for each ALT allele, in the same order as listed">
AF Fraction of evidence in support of alternate allele(s)">
SVTYPE Type of structural variant">
SVLEN Difference in length between REF and ALT alleles">
END End position of the variant described in this record">
IMPRECISE Imprecise change from local reassembly (ALT contains Ns)

Take-homes

- Error mode of sequencing technology is key to alignment.
- Post-processing is key:
 - Duplicate marking
 - INDEL realignment